

Video Transcription

- **Primary Model: faster-whisper** with the large-v3 model.

Whisper models are **large pre-trained models**. Training a model from scratch to understand human speech requires massive amounts of data and computing power. Using a pre-trained model like Whisper allows project to leverage its existing knowledge of dozens of languages and accents instantly, without any training needed. I'm here **fine-tuning the output**, not the model itself.

- faster-whisper is a highly optimized version of OpenAI's Whisper. It's much faster and uses less memory because it's built on CTranslate2. The large-v3 model I chose because it is the most accurate for multilingual transcription and translation.
- On newer Python versions or if my GPU runs out of memory, it falls back to the standard openai-whisper library with a medium.en model, which is a good balance of speed and accuracy.

Key Algorithms & Techniques

- **Audio Pre-processing (FFmpeg, Pydub, noisereduce)**: Used before sending audio to the model. It **converts any audio/video file to a standardized 16kHz mono WAV format**, normalizes its volume, and reduces background noise to improve transcription accuracy.
- **Automatic Speech Recognition (ASR)**: The core task performed by the Whisper model. It **converts audio signals into text**.
- **Voice Activity Detection (VAD - Silero VAD)**: Used during transcription to identify and only transcribe parts of the audio where someone is speaking, **ignoring silent gaps and reducing errors**. This is a key feature of faster-whisper.
- **Beam Search**: Used during decoding (the model's text generation step). It keeps several likely options (beam_size=5) open at each step, leading to more accurate and coherent transcriptions than simpler greedy methods.
- **Text Post-processing (Custom Algorithms)**: Used after transcription to clean the raw output. This includes:
 - **Jaccard Similarity**: To detect and remove repetitive sentences.
 - **Rule-based Filters**: To collapse repeated words and trim nonsensical endings (e.g., "आपको आपको आपको").
 - **Context-aware Corrections**: To fix common informal shorthand (e.g., "u" -> "you").

The Main Steps (Pipeline)

1. **Input**: User provides a video/audio file or a YouTube URL.
2. **Audio Extraction**: FFmpeg extracts the audio track.
3. **Pre-processing**: Audio is converted to 16kHz mono, normalized, and cleaned of noise.
4. **Transcription**: The pre-processed audio is fed to the Whisper model, which uses VAD and Beam Search to generate raw text with word-level timestamps.
5. **Post-processing**: Custom algorithms clean the text, remove repetitions, and improve formatting.
6. **Output**: The final text is saved, along with an SRT file (for subtitles) and a JSON file (with precise word timestamps).



